

Machine Learning 1

Lecture 07 - Logistic Regression - Neural Nets

Patrick Forré

1 Logistic Regression

2 Supervised Learning: Neural Nets

Linear Classification: Logistic Regression for two classes (I)

- Given: Data set $D = (x_1, \dots, x_N)^T$ with binary classes $T = (t_1, \dots, t_N)^T$ with $t_i \in \{c_0, c_1\} = \{0, 1\}$.
- Basis functions: $\phi = \phi(x) = (\phi_0(x), \dots, \phi_M(x))^T$ with $\phi_0 \equiv 1$.
- Model assumption of Logistic Regression: The posterior probability $p(c_1|\phi)$ is the sigmoid of a linear function in the feature vector ϕ :

$$p(c_1|\phi, w) = \sigma(w^T \phi)$$

with weight vector $w = (w_0, \dots, w_M) \in \mathbb{R}^{M+1}$.

- Conditional distribution:

$$\begin{aligned} p(t|\phi, w) &= \begin{cases} \sigma(w^T \phi) & \text{if } t = 1, \\ 1 - \sigma(w^T \phi) & \text{if } t = 0, \end{cases} \\ &= \sigma(w^T \phi)^t \cdot (1 - \sigma(w^T \phi))^{1-t}. \end{aligned}$$

Linear Classification: Logistic Regression for two classes (II)

- Conditional Likelihood (under i.i.d. assumptions):

$$\begin{aligned} p(T|\Phi, w) &= \prod_{n=1}^N p(t_n|\phi(x_n), w) \\ &= \prod_{n=1}^N \sigma(w^T \phi(x_n))^{t_n} \cdot (1 - \sigma(w^T \phi(x_n)))^{1-t_n} \\ &= \prod_{n=1}^N y_n^{t_n} \cdot (1 - y_n)^{1-t_n} \end{aligned}$$

with $y_n := \sigma(w^T \phi(x_n))$. Put $Y = (y_1, \dots, y_N)^T$.

- For the maximum likelihood approach we either needed to know $p(\Phi|w)$ or at least assume that it does not depend on w .
- This leads to maximizing the conditional likelihood w.r.t. w .
- This is equivalent to minimizing the cross-entropy error:

$$\begin{aligned} E(w) &= -\ln p(T|\Phi, w) \\ &= -\sum_{n=1}^N [t_n \ln y_n + (1 - t_n) \ln(1 - y_n)] \end{aligned}$$

- $E(w)$ is convex, but no closed form solution exists (due to the non-linearity of σ).
- For minimizing $E(w)$ use e.g. stochastic gradient descent or iteratively reweighted least squares.

Theorem (Chain Rule of Calculus)

For two differentiable and composable functions f, g we have:

$$(f \circ g)'(x) = f'(g(x)) \cdot g'(x).$$

For three functions h, f, g we have:

$$(h \circ f \circ g)'(x) = h'(f(g(x))) \cdot f'(g(x)) \cdot g'(x).$$

Example: Logistic sigmoid function $\sigma(x) = (1 + \exp(-x))^{-1}$.

$$\begin{aligned}\sigma'(x) &= (-1)(1 + \exp(-x))^{-2} \cdot \exp(-x) \cdot (-1) \\ &= \frac{1}{1 + \exp(-x)} \frac{\exp(-x)}{1 + \exp(-x)} \\ &= \frac{1}{1 + \exp(-x)} \frac{1}{\exp(x) + 1} \\ &= \sigma(x)\sigma(-x) \\ &= \sigma(x)(1 - \sigma(x)).\end{aligned}$$

Stochastic Gradient Descent for Cross-Entropy (I)

- Given: Data set $D = (x_1, \dots, x_N)^T$ with binary classes $T = (t_1, \dots, t_N)^T$ with $t_i \in \{c_0, c_1\} = \{0, 1\}$.
- Put $\Phi = (\phi(x_1), \dots, \phi(x_N))^T$ with basis functions $\phi(x) = (\phi_0(x), \dots, \phi_M(x))^T$ and $y_n = \sigma(w^T \phi(x_n))$.
- Cross entropy error for one observation:
 $E_n(w) = -(t_n \ln y_n + (1 - t_n) \ln(1 - y_n))$.

$$\begin{aligned}\frac{\partial \ln \sigma(w^T \phi(x_n))}{\partial w_j} &= \frac{1}{\sigma(w^T \phi(x_n))} \cdot \sigma'(w^T \phi(x_n)) \cdot \frac{w^T \phi(x_n)}{\partial w_j} \\ &= \frac{\sigma(w^T \phi(x_n)) \sigma(-w^T \phi(x_n))}{\sigma(w^T \phi(x_n))} \cdot \frac{\sum_m w_m \phi_m(x_n)}{\partial w_j} \\ &= \sigma(-w^T \phi(x_n)) \cdot \phi_j(x_n).\end{aligned}$$

$$\frac{\partial^2 \ln \sigma(w^T \phi(x_n))}{\partial^2 w_j} = -\sigma(w^T \phi(x_n)) \cdot \sigma(-w^T \phi(x_n)) \cdot \phi_j(x_n)^2 < 0.$$

$$\begin{aligned}\frac{\partial E_n(w)}{\partial w_j} &= -t_n \frac{\partial \ln y_n}{\partial w_j} - (1 - t_n) \frac{\partial \ln(1 - y_n)}{\partial w_j} \\ &= -t_n \frac{\partial \ln \sigma(w^T \phi(x_n))}{\partial w_j} - (1 - t_n) \frac{\partial \ln(1 - \sigma(w^T \phi(x_n)))}{\partial w_j} \\ &= -t_n \sigma(-w^T \phi(x_n)) \cdot \phi_j(x_n) + (1 - t_n) \sigma(w^T \phi(x_n)) \cdot \phi_j(x_n) \\ &= (-t_n(1 - y_n) + (1 - t_n)y_n) \cdot \phi_j(x_n) \\ &= (y_n - t_n) \cdot \phi_j(x_n).\end{aligned}$$

Stochastic Gradient Descent for Cross-Entropy (II)

- Cross-entropy error for one observation:
 $E_n(w) = -(t_n \ln y_n + (1 - t_n) \ln(1 - y_n))$ with
 $y_n = \sigma(w^T \phi(x_n))$.
- $\frac{\partial E_n(w)}{\partial w_j} = (y_n - t_n) \cdot \phi_j(x_n)$.
- Gradient: $\nabla E_n(w) = (\frac{\partial E_n(w)}{\partial w_0}, \dots, \frac{\partial E_n(w)}{\partial w_M})^T = (y_n - t_n) \cdot \phi(x_n)$.
- Cross-entropy: $E(w) = \sum_{n=1}^N E_n(w) = -\sum_{n=1}^N [t_n \ln y_n + (1 - t_n) \ln(1 - y_n)]$.
- Stochastic Gradient Descent:
 - Carefully choose an initial value: $w^{(0)}$ (often $w^{(0)} = 0$ works).
 - Carefully choose a learning rate $\eta > 0$.
 - Randomly choose an observation x_n and use update rule:
$$\begin{aligned}w^{(\tau+1)} &:= w^{(\tau)} - \eta \nabla E(w^{(\tau)}) \\ &= w^{(\tau)} - \eta \cdot (y_n^{(\tau)} - t_n) \cdot \phi(x_n)\end{aligned}$$
 - Iterating the last point will make the algorithm converges (if η is not too big): $E(w)$ is convex (Hessian positive definite).
 - If η is too small the algorithm is too slow.
- We end up with an approximate minimizer w^* of $E(w)$.

Compare: Stochastic Gradient Descent for Linear Regression

- Sum-of-Squares error: $E(w) = \sum_{n=1}^N E_n(w)$ with:
- Sum-of-Squares error for one observation:
 $E_n(w) = \frac{1}{2}(y_n - t_n)^2$ with $y_n = w^T \phi(x_n)$.
- $\frac{\partial E_n(w)}{\partial w_j} = (y_n - t_n) \cdot \phi_j(x_n)$.
- Gradient: $\nabla E_n(w) = (\frac{\partial E_n(w)}{\partial w_0}, \dots, \frac{\partial E_n(w)}{\partial w_M})^T = (y_n - t_n) \cdot \phi(x_n)$.
- Stochastic Gradient Descent with same update rule as in Logistic Regression:

$$\begin{aligned} w^{(\tau+1)} &:= w^{(\tau)} - \eta \nabla E(w^{(\tau)}) \\ &= w^{(\tau)} - \eta \cdot (y_n^{(\tau)} - t_n) \cdot \phi(x_n) \end{aligned}$$

Iteratively Reweighted Least Squares for Cross-Entropy

- Given: Data set $D = (x_1, \dots, x_N)^T$ with binary classes $T = (t_1, \dots, t_N)^T$ with $t_i \in \{c_0, c_1\} = \{0, 1\}$.
- Put $\Phi = (\phi(x_1), \dots, \phi(x_N))^T$ with basis functions $\phi(x) = (\phi_0(x), \dots, \phi_M(x))^T$.
- Put $y_n = \sigma(w^T \phi(x_n))$ and $Y = (y_1, \dots, y_N)^T$ with $w \in \mathbb{R}^{M+1}$.
- $E(w) = -\sum_{n=1}^N [t_n \ln y_n + (1 - t_n) \ln(1 - y_n)]$.
- Gradient: $\nabla E(w) = \sum_{n=1}^N (y_n - t_n) \phi(x_n) = \Phi^T (Y - T)$.
- Hessian: $H(w) = \sum_{n=1}^N y_n(1 - y_n) \phi(x_n) \phi(x_n)^T = \Phi^T R \Phi$, with
- diagonal matrix R with entries $R_{nn} = y_n(1 - y_n)$.
- Newton-Raphson update:

$$\begin{aligned} w^{(t+1)} &:= w^{(t)} - H(w^{(t)})^{-1} \nabla E(w^{(t)}) \\ &= w^{(t)} - (\Phi^T R^{(t)} \Phi)^{-1} \Phi^T (Y^{(t)} - T) \\ &= (\Phi^T R^{(t)} \Phi)^{-1} [\Phi^T R^{(t)} \Phi w^{(t)} - \Phi^T (Y^{(t)} - T)] \\ &= (\Phi^T R^{(t)} \Phi)^{-1} \Phi^T R^{(t)} Z^{(t)}, \\ Z^{(t)} &= \Phi w^{(t)} - (R^{(t)})^{-1} (Y^{(t)} - T) \end{aligned}$$

- We end up with an approximate minimizer w^* of $E(w)$.

Geometry of Gradient Descent and Newton Optimization

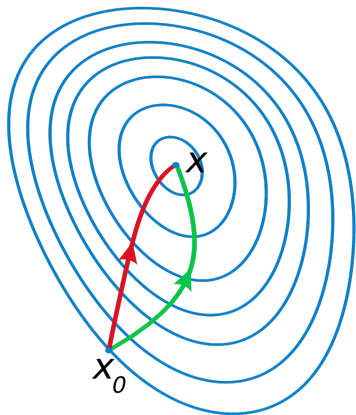


Figure: Contours of a convex error function $E(w)$. $w = X$ is the minimum of $E(w)$ and $w = X_0$ a starting point. Green: Gradient descent follows the steepest descent at each point, orthogonal to the contours. Red: Newton-Raphson method also takes the curvature into account to shorten the way. (Source: Wikipedia - Newton's method in optimization)

Classification with Logistic Regression

- Given: Data set $D = (x_1, \dots, x_N)^T$ with binary classes $T = (t_1, \dots, t_N)^T$ with $t_i \in \{c_0, c_1\} = \{0, 1\}$.
- Basis functions: $\phi = \phi(x) = (\phi_0(x), \dots, \phi_M(x))^T$.
- Model assumption of Logistic Regression:
 $p(c_1|\phi, w) = \sigma(w^T \phi)$.
- Minimizing the cross-entropy error:

$$\begin{aligned} E(w) &= -\ln p(T|\Phi, w) \\ &= -\sum_{n=1}^N [t_n \ln y_n + (1 - t_n) \ln(1 - y_n)]. \end{aligned}$$

- Using stochastic gradient descent or iterative reweighted least squares we end up with a approximate minimizer w^* of $E(w)$.
- We assign a new data point x to class c_1 if $\sigma((w^*)^T \phi(x)) > \frac{1}{2}$, i.e. if $(w^*)^T \phi(x) > 0$.
- Decision regions: $\mathcal{R}_1 = \{x | (w^*)^T \phi(x) > 0\}$ and $\mathcal{R}_0 = \{x | (w^*)^T \phi(x) < 0\}$.
- Decision boundaries: $\mathcal{B} = \{x | (w^*)^T \phi(x) = 0\}$.

Logistic Regression for multiple classes

- Data $D = (x_1, \dots, x_N)^T$ with $T = (t_1, \dots, t_N)^T$ of K -dim one-vs-the-rest vectors $t_i = (0, \dots, 1, \dots, 0)^T$.
- Model assumption of Logistic Regression:

$$p(c_k | \phi, w_1, \dots, w_K) = \sigma_k(w_1^T \phi, \dots, w_K^T \phi),$$

with weight vectors $w_k = (w_{k,0}, \dots, w_{k,M}) \in \mathbb{R}^{M+1}$.

- Put $y_{nk} := \sigma_k(w_1^T \phi(x_n), \dots, w_K^T \phi(x_n))$.
- Minimize the cross-entropy error w.r.t. w :

$$E(W) = -\ln p(T | \Phi, W) = -\sum_{n=1}^N \sum_{k=1}^K t_{nk} \ln y_{nk}.$$

- Gradient: $\nabla_{w_j} E(W) = \sum_{n=1}^N (y_{nj} - t_{nj}) \phi(x_n)$
- Hessian: $\nabla_{w_k} \nabla_{w_j} E(W) = -\sum_{n=1}^N y_{nk} (\mathbb{1}_{nj} - y_{nj}) \phi(x_n) \phi(x_n)^T$.
- We assign x to class c_k if $\sigma_k > \sigma_j$ for all $j \neq k$, i.e.:
- Decision regions: $\mathcal{R}_k = \{x | (w_k^*)^T \phi(x) > (w_j^*)^T \phi(x), \forall j \neq k\}$.
- Decision boundaries: $\mathcal{B}_{jk} = \{x | (w_j^*)^T \phi(x) = (w_k^*)^T \phi(x)\}$.

Bayesian Model Comparison and BIC

- Given: model \mathcal{M}_i with parameters \mathcal{W}_i of dim M_i and data D .
- Bayesian Model Comparison: Highest model evidence:

$$p(D|\mathcal{M}_i) = \int_{\mathcal{W}_i} p(D|w, \mathcal{M}_i) p(w|\mathcal{M}_i) dw.$$

- (In model) Posterior: $p(w|D, \mathcal{M}_i) = \frac{p(D|w, \mathcal{M}_i)p(w|\mathcal{M}_i)}{p(D|\mathcal{M}_i)}$.
- Laplace approximation (see next slide) of maximum a posteriori for model \mathcal{M}_i :

$$\begin{aligned} p(w|D, \mathcal{M}_i) &\approx \mathcal{N}(w|w_{\text{MAP}}, A^{-1}), \\ A &= -\nabla\nabla \ln p(w_{\text{MAP}}|D, \mathcal{M}_i) \\ &= -\nabla\nabla \ln p(D|w_{\text{MAP}}, \mathcal{M}_i)p(w_{\text{MAP}}|\mathcal{M}_i). \end{aligned}$$

- This leads to the approximation: $\ln p(D|\mathcal{M}_i) \approx \ln p(D|w_{\text{MAP}}, \mathcal{M}_i) + \ln p(w|\mathcal{M}_i) + \frac{M_i}{2} \ln(2\pi) - \frac{1}{2} \ln |A|$.
- Broad Gaussian prior and A of full rank gives us BIC:

$$\ln p(D|\mathcal{M}_i) \approx \ln p(D|w_{\text{MAP}}, \mathcal{M}_i) - \frac{M_i}{2} \ln N + \text{const.}$$

- This is a very rough approximation and the assumptions often fail in practice.

Laplace Approximation

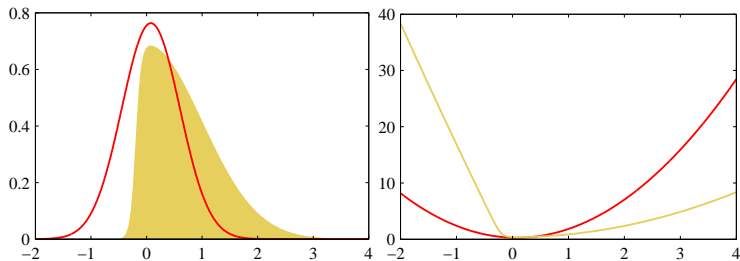


Figure: Left: Distribution $p(z)$ approximated by Gaussian distribution $\mathcal{N}(z|z_0, A^{-1})$ around mode $z = z_0$ with Hessian matrix $A = -\nabla\nabla \ln p(z)|_{z=z_0}$. Right: $-\ln p(x)$ (yellow) and $-\ln \mathcal{N}(z|z_0, A^{-1})$ (red). (Bishop 4.14)

Bayesian Logistic Regression

- Likelihood: $p(T|\Phi, w) = \prod_{n=1}^N y_n^{t_n} \cdot (1 - y_n)^{1-t_n}$ with $y_n := \sigma(w^T \phi(x_n))$.
- Gaussian prior: $p(w) = \mathcal{N}(w|\mu_0, \Sigma_0)$.
- Posterior then is:

$$\begin{aligned} \ln p(w|\Phi, T) &= -\frac{1}{2}(x - \mu_0)^T \Sigma_0^{-1}(x - \mu_0) \\ &\quad + \sum_n [t_n \ln y_n + (1 - t_n) \ln(1 - y_n)] + \text{const.} \end{aligned}$$

- Laplace approximation of the posterior:

$$\begin{aligned} p(w|\Phi, T) &\approx \mathcal{N}(w|w_{\text{MAP}}, \Sigma_N) \\ \Sigma_N^{-1} &= -\nabla \nabla \ln p(w|T, \Phi) \\ &= \Sigma_0^{-1} + \sum_{n=1}^N y_n(1 - y_n) \phi(x_n) \phi(x_n)^T. \end{aligned}$$

- Predictive distribution: $p(c_1|\phi, \Phi, T) \approx$

$$\int \sigma(w^T \phi) \mathcal{N}(w|w_{\text{MAP}}, \Sigma_N) dw \approx \sigma \left(\left(1 + \frac{\pi}{8} \Phi^T \Sigma_N \Phi\right)^{-1/2} w_{\text{MAP}}^T \phi \right).$$

1 Logistic Regression

2 Supervised Learning: Neural Nets

Neural Nets: The Idea

- First consider Linear Regression and Logistic Regression.
- Training data: D with targets $T = (t_1, \dots, t_N)^T$.
- Targets (simplest case): LinReg: $t_i \in \mathbb{R}$. LogReg: $t_i \in \{0, 1\}$.
- Features: $\phi(x) = (\phi_0(x), \dots, \phi_M(x))^T$ with $\phi_0 \equiv 1$.
- Model functions in LinReg: $y(x, w) = w^T \phi(x)$.
- Model functions in LogReg: $y(x, w) = \sigma(w^T \phi(x))$.
- Problem: Feature functions need to be known or handcrafted!
- Idea: Create "flexible" non-linear features and learn them:

$$\phi_m(x) = \sigma(w'_m{}^T x) = \sigma \left(\sum_{d=0}^D w'_{m,d} \cdot x_d \right).$$

with weights $w'_m \in \mathbb{R}^{D+1}$, $x = (1, x_1, \dots, x_D)$ (here: components, not observations), $x_0 \equiv 1$.

- Regression: $y(x, w, w') = \sum_{m=0}^M w_m \cdot \sigma \left(\sum_{d=0}^D w'_{m,d} \cdot x_d \right)$.
- Classification:
 $y(x, w, w') = \sigma \left(\sum_{m=0}^M w_m \cdot \sigma \left(\sum_{d=0}^D w'_{m,d} \cdot x_d \right) \right)$.
- Caution: The inner functions are not linear anymore!!!

Question (Why should this work?)

Why should the approach to model functions like:

$$y(x, w, w') = \sum_{m=0}^M w_m \cdot \sigma \left(\sum_{d=0}^D w'_{m,d} \cdot x_d \right),$$

work? Why the logistic sigmoid function σ ?

Theorem (Universal Approximator)

Let f be any continuous function on a compact area of \mathbb{R}^D and σ any fixed analytic function which is not a polynomial (e.g. σ logistic function). Given any small number $\epsilon > 0$ of an acceptable error we can find a number M and weights $w_m, w'_{m,d} \in \mathbb{R}$ such that:

$$|f(x) - y(x, w, w')| < \epsilon.$$

Caution: *With smaller ϵ we usually have to take bigger M !*

Example: Approximation Power of Neural Nets

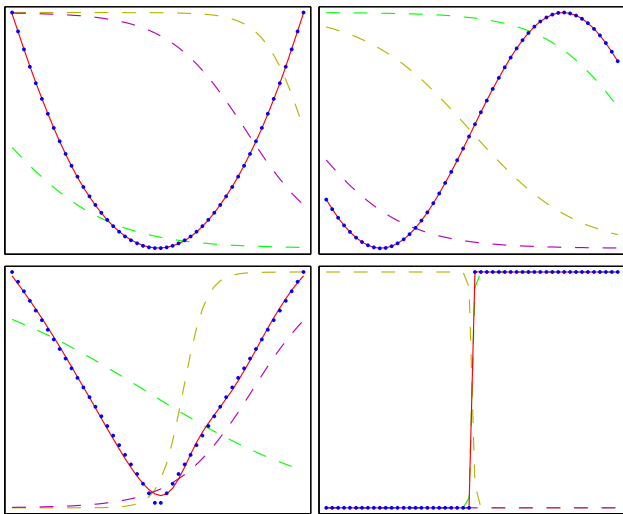


Figure: Approximation power of neural net with $M = 3$ features. $N = 50$ data points. a) $f(x) = x^2$, b) $f(x) = \sin(x)$, c) $f(x) = |x|$, d) $f(x) =$ step function. Dashed lines are the curves of the features. (Bishop 5.3)

Example: Classification Power of Neural Nets

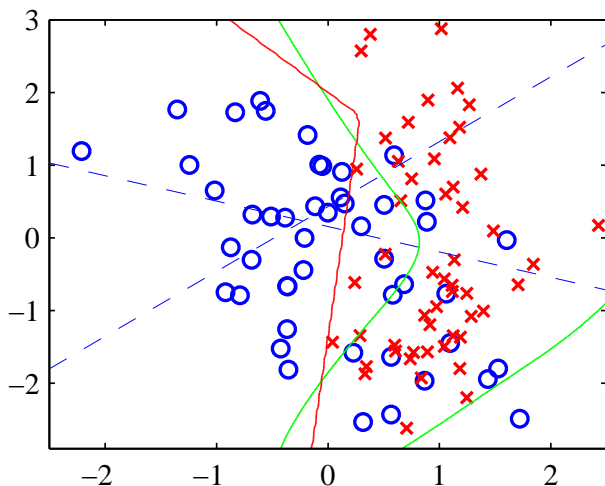


Figure: Classification power of neural net with $M = 2$ features. Decision boundary in red. Optimal boundary in green. Dashed blue lines are $z = 0.5$ contours of the features. (Bishop 5.4)

Neural Nets: Multiple Outputs and Multiple Classes

- Training data: $D = (x_1, \dots, x_N)^T$ with $T = (t_1, \dots, t_N)^T$.
- Regression: $t_i \in \mathbb{R}^K$ (K -components).
- Classification: K -dim one-vs-the-rest: $t_i = (0, \dots, 1, \dots, 0)^T$.
- Model functions of Neural Nets for Regression:

$$y_k(x, w^{(1)}, w^{(2)}) = \sum_{m=0}^M w_{k,m}^{(2)} \cdot \sigma \left(\sum_{d=0}^D w_{m,d}^{(1)} \cdot x_d \right),$$

where σ is the logistic sigmoid function.

After learning $w^{(1)}, w^{(2)}$ the target $t = (t_1, \dots, t_K)^T$ of new data x is predicted by $(y_1(x), \dots, y_K(x))^T$.

- Model functions of Neural Nets for Classification:

$$y_k(x, w^{(1)}, w^{(2)}) = \sigma_k \left(\left(\sum_{m=0}^M w_{k',m}^{(2)} \cdot \sigma \left(\sum_{d=0}^D w_{m,d}^{(1)} \cdot x_d \right) \right)_{k'=1,\dots,K} \right)$$

where σ_k is k -th comp. of softmax or the sigmoid of k -comp..

After learning $w^{(1)}, w^{(2)}$ a new data point x is assigned to class c_k if $y_k(x) > y_j(x)$ for all $j \neq k$.

Graphical Representation of Neural Nets

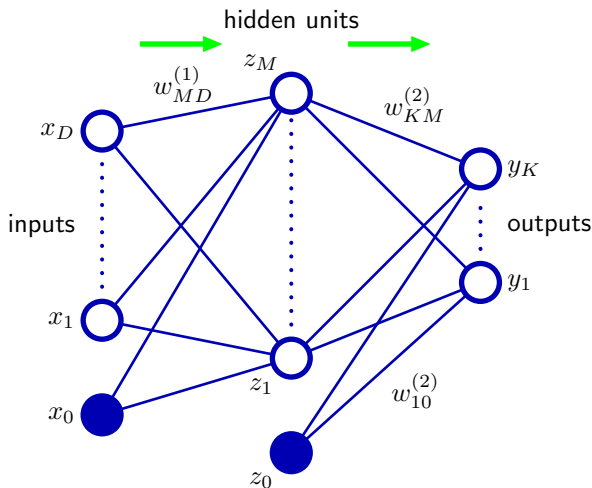


Figure: Edges correspond to weights $w_{i,j}^{(\ell)}$ and nodes to input, output and hidden units $z_m = \sigma(\sum_d w_{m,d}^{(1)} x_d)$. Big blue nodes correspond to constant functions $x_0 = 1, z_0 = 1$ and their edges to bias terms. (Bishop 5.1)